

```
import numpy
import urllib
import scipy.optimize
import random
from sklearn import svm
import os
import yaml
import random
import operator
```

```
def parseData(fname):
    for l in urllib.request.urlopen(fname):
        yield eval(l)
data = list(parseData("http://jmcauley.ucsd.edu/cse190/data/beer/beer_50000.json"))
```

#Question 1

```
print('QUESTION 1')
```

```
ratings = {'1-star':0, '2-star':0, '3-star':0, '4-star':0, '5-star':0}
```

```
for d in data:
```

```
    ratings[str(int(d['review/taste']))+'-star'] += 1
```

```
print('Rounded down ratings: {}'.format(ratings))
```

#Question 2

```
print('QUESTION 2')
listOfBeers = set()
ratings = dict()
averageRatings = dict()
for d in data:
    listOfBeers.add(d['beer/name'])
for beerName in listOfBeers:
    ratings[beerName] = 0
    averageRatings[beerName] = 0
for d in data:
    ratings[d['beer/name']] += 1
    averageRatings[d['beer/name']] += d['review/taste']
for key in ratings.keys():
    averageRatings[key] = -1 if ratings[key] < 5 else averageRatings[key]/ratings[key]
maximum = max(averageRatings.values())
result = []
for item in averageRatings.items():
    if item[1] == maximum:
        result.append(item)
print(result)
```

#Question 3

```
print('QUESTION 3')
data2 = [d for d in data if 'beer/ABV' in d.keys() and 'beer/style' in d.keys()]
def feature(datum):
    feat = [1]
    if datum['beer/style'] == 'Hefeweizen':
        feat.append(1)
    else:
        feat.append(0)
    feat.append(datum['beer/ABV'])
    return feat
X = [feature(d) for d in data2]
y = [d['review/taste'] for d in data2]
theta,residuals,rank,s = numpy.linalg.lstsq(X, y, rcond=None)
print(theta)
#theta[0] shows us the offset. Theta[1] is the weighting of the beer being a
#Hefeweizen on the taste score, in this case it has negative impact on the score.
#Theta[2] shows the impact of the beer alcohol content on the taste score, in this case
#higher is better
```

#Question 4

```
print('QUESTION 4')
train = data2[:len(data2)//2]
X_train = [feature(d) for d in train]
y_train = [d['review/taste'] for d in train]
theta,residuals,ranks = numpy.linalg.lstsq(X_train, y_train, rcond=-1)
predictions = numpy.matmul(theta, numpy.transpose(X_train))
print('Error for train data =', numpy.sum(numpy.square(numpy.subtract(predictions, y_train))))

test = data2[len(data2)//2:]
X_test = [feature(d) for d in test]
y_test = [d['review/taste'] for d in test]
predictions = numpy.matmul(theta, numpy.transpose(X_test))
print('Error for test data =', numpy.sum(numpy.square(numpy.subtract(predictions, y_test))))
```

#Question 5

```
print('QUESTION 5')
dataCopy = data2.copy()
for i in range(10):
    random.shuffle(dataCopy)

train = dataCopy[:len(dataCopy)//2]
X_train = [feature(d) for d in train]
y_train = [d['review/taste'] for d in train]
theta,residuals,rank,s = numpy.linalg.lstsq(X_train, y_train, rcond=-1)
predictions = numpy.matmul(theta, numpy.transpose(X_train))
print('Error for train data =', numpy.sum(numpy.square(numpy.subtract(predictions, y_train))))

test = dataCopy[len(dataCopy)//2:]
X_test = [feature(d) for d in test]
y_test = [d['review/taste'] for d in test]
predictions = numpy.matmul(theta, numpy.transpose(X_test))
print('Error for test data =', numpy.sum(numpy.square(numpy.subtract(predictions, y_test))))

#The data that is pulled from the website might be in some certain order and dividing it before
shuffling
#might result in skewed data
```

#Question 6

```
print('QUESTION 6')
def feature(datum):
    #feat = [1]
    feat = []
    feat.append(datum['review/taste'])
    feat.append(datum['review/appearance'])
    feat.append(datum['review/aroma'])
    feat.append(datum['review/palate'])
    feat.append(datum['review/overall'])
    return feat

dataCopy = data2.copy()
for i in range(10):
    random.shuffle(dataCopy)

train = dataCopy[:len(dataCopy)//2]
X_train = [feature(d) for d in train]
y_train = [1 if d['beer/style'] == 'Hefeweizen' else 0 for d in train]
clf = svm.SVC(C=1000, kernel='rbf', gamma='scale', degree=2)
clf.fit(X_train, y_train)
train_predictions = clf.predict(X_train)
print('Accuracy for train data =', 100-
(100*numpy.sum(numpy.square(numpy.subtract(train_predictions, y_train)))/len(y_train)), '%')

test = dataCopy[len(dataCopy)//2:]
X_test = [feature(d) for d in test]
y_test = [1 if d['beer/style'] == 'Hefeweizen' else 0 for d in test]
test_predictions = clf.predict(X_test)
print('Accuracy for test data =', 100-
(100*numpy.sum(numpy.square(numpy.subtract(test_predictions, y_test)))/len(y_test)), '%')
```

#Question 7

```
print('QUESTION 7')
def feature(datum):
    feat = [1]
    feat.append(datum['beer/ABV'])
    feat.append(1 if 'malt' in datum['review/text'] or 'malty' in datum['review/text'] else 0)
    feat.append(datum['review/timeUnix'])
    return feat

train = dataCopy[:len(dataCopy)//2]
X_train = [feature(d) for d in train]
y_train = [1 if d['beer/style'] == 'Hefeweizen' else 0 for d in train]
clf2 = svm.SVC(C=1000, kernel='rbf', gamma='scale')
clf2.fit(X_train, y_train)
train_predictions = clf2.predict(X_train)
print('Accuracy for train data =', 100-
(100*numpy.sum(numpy.square(numpy.subtract(train_predictions, y_train)))/len(y_train)), '%')

test = dataCopy[len(dataCopy)//2:]
X_test = [feature(d) for d in test]
y_test = [1 if d['beer/style'] == 'Hefeweizen' else 0 for d in test]
test_predictions = clf2.predict(X_test)
print('Accuracy for test data =', 100-
(100*numpy.sum(numpy.square(numpy.subtract(test_predictions, y_test)))/len(y_test)), '%')

#The features used are:
#[1, beer/ABV, Boolean indicating if review text contains the word 'malt' or 'malty',
review/timeUnix]
```


#Question 8

```
print('QUESTION 8')
```

```
c = [0.1, 10, 1000, 100000]
```

```
for reg in c:
```

```
    clf = svm.SVC(C=reg, kernel='rbf', gamma='scale')
```

```
    clf.fit(X_train, y_train)
```

```
    train_predictions = clf.predict(X_train)
```

```
    print('-----For regularization factor =', reg, '-----')
```

```
    print('Accuracy for train data =', 100-
```

```
(100*numpy.sum(numpy.square(numpy.subtract(train_predictions, y_train)))/len(y_train)), '%')
```

```
    test_predictions = clf.predict(X_test)
```

```
    print('Accuracy for test data =', 100-
```

```
(100*numpy.sum(numpy.square(numpy.subtract(test_predictions, y_test)))/len(y_test)), '%')
```

Printed Results

QUESTION 1

Rounded down ratings: {'1-star': 554, '2-star': 2723, '3-star': 12934, '4-star': 29458, '5-star': 4331}

QUESTION 2

[('Founders CBS Imperial Stout', 4.6970172684458396)]

QUESTION 3

[3.11795084 -0.05637406 0.10877902]

QUESTION 4

Error for train data = 12099.201400335607

Error for test data = 10592.663029965459

QUESTION 5

Error for train data = 11435.60803063151

Error for test data = 11053.513210952766

QUESTION 6

Accuracy for train data = 98.776 %

Accuracy for test data = 98.712 %

QUESTION 7

Accuracy for train data = 99.172 %

Accuracy for test data = 97.968 %

QUESTION 8

-----For regularization factor = 0.1 -----

Accuracy for train data = 98.748 %

Accuracy for test data = 98.78 %

-----For regularization factor = 10 -----

Accuracy for train data = 98.872 %

Accuracy for test data = 98.56 %

-----For regularization factor = 1000 -----

Accuracy for train data = 99.172 %

Accuracy for test data = 97.968 %

-----For regularization factor = 100000 -----

Accuracy for train data = 99.592 %

Accuracy for test data = 97.508 %