

```
In [2]: import numpy as np
import urllib
import scipy.optimize
import random
from collections import defaultdict
import nltk
import string
from nltk.stem.porter import *
from sklearn import linear_model
import operator
from math import log
import re
from random import shuffle
from sklearn.metrics import mean_squared_error
```

```
In [3]: def parseData(fname):
        for l in urllib.request.urlopen(fname):
            yield eval(l)

        ### Just the first 5000 reviews

        print ("Reading data...")
        fullData = list(parseData("http://jmcauley.ucsd.edu/cse190/data/beer/beer_50000.json"))
        data = fullData[:5000]
        shuffle(fullData)
        trainData = fullData[:5000]
        valData = fullData[5000:10000]
        testData = fullData[10000:]
        print ("done")
        translator = str.maketrans(dict.fromkeys(string.punctuation))

Reading data...
done
```

In [10]: *### How many unique words are there?*

```
wordCount = defaultdict(int)
for d in data:
    for w in d['review/text'].split():
        wordCount[w] += 1

print(len(wordCount))

### Ignore capitalization and remove punctuation

wordCount = defaultdict(int)
punctuation = set(string.punctuation)
stemmer = PorterStemmer()
for d in data:
    r = ''.join([c for c in d['review/text'].lower() if not c in punctuation])
    for w in r.split():
        #w = stemmer.stem(w) # with stemming
        wordCount[w] += 1

### Just take the most popular words...

wordCount = defaultdict(int)
punctuation = set(string.punctuation)
for d in data:
    r = ''.join([c for c in d['review/text'].lower() if not c in punctuation])
    for w in r.split():
        wordCount[w] += 1

counts = [(wordCount[w], w) for w in wordCount]
counts.sort()
counts.reverse()

words = [x[1] for x in counts[:1000]]
toBeFound = words.copy()
```

36225

```

In [11]: ### Bigram counts
bigramCount = defaultdict(int)
reviewBigramContent = []
for d in data:
    text = d['review/text']
    removed = text.translate(translator)
    lowered = removed.lower()
    wordList = lowered.split()
    prev = None
    reviewSet = []
    for w in wordList:
        if (prev is None):
            bigramCount[('/',w)] += 1
            reviewSet.append('/',w)
        else:
            bigramCount[(prev,w)] += 1
            reviewSet.append((prev,w))
        prev = w
    reviewBigramContent.append(reviewSet)
frequentWords = sorted(list(bigramCount.items()), key=operator.itemgetter(1))
frequentWords.reverse()
bigrams = set(bigramCount.keys())
bigramNoDup = list(bigrams)
print('')
print('Total number of unique bigrams: {}'.format(len(frequentWords)))
print('Top 5 bigrams:')
for i in range(5):
    print('Bigram {} with count {}'.format(frequentWords[i][0], frequentWords[i][1]))

wordId = dict(zip(words, range(len(words))))
wordSet = set(words)

def feature(datum, wordKey):
    feat = [0]*len(words)
    r = ''.join([c for c in datum['review/text'].lower() if not c in punctuation])
    for w in r.split():
        if w in words:
            feat[wordKey[w]] += 1
    feat.append(1) #offset
    return feat

X = [feature(d,wordId) for d in data]
y = [d['review/overall'] for d in data]

#No regularization


#With regularization
clf = linear_model.Ridge(1.0, fit_intercept=False)
clf.fit(X, y)
theta = clf.coef_

```

```

predictions = clf.predict(X)
err = mean_squared_error(y,predictions)
print('\nMSE of unigram model: {}'.format(err))

### Bigram model
def featureBi(bigramList,wordKey):
    feat = [0] * len(bigrams)
    for w in bigramList:
        feat[wordKey[w]] += 1
    feat.append(1)
    return feat
wordId = dict(zip(bigramNoDup, range(len(bigramNoDup))))
X = [featureBi(d,wordId) for d in reviewBigramContent]
y = [d['review/overall'] for d in data]
clf = linear_model.Ridge(1.0, fit_intercept=False)
clf.fit(X,y)
theta = clf.coef_
predictions = clf.predict(X)
err = mean_squared_error(y,predictions)
print('\nMSE of bigram model: {}'.format(err))

```

Total number of unique bigrams: 182902

Top 5 bigrams:

Bigram ('with', 'a') with count 4587

Bigram ('in', 'the') with count 2595

Bigram ('of', 'the') with count 2245

Bigram ('is', 'a') with count 2056

Bigram ('on', 'the') with count 2033

MSE of unigram model: 0.2787546353022137

MSE of bigram model: 0.0004163648437118055

```

In [12]: reviewsAndWords = []
allWords = set()
for d in data:
    text = d['review/text']
    words = ((text.translate(translator)).lower()).split()
    reviewsAndWords.append(words)
    for w in words:
        allWords.add(w)

```

```

In [13]: checkWords = ['foam','smell','banana','lactic','tart']
totalDocuments = len(reviewsAndWords)
inverseFrequencies = defaultdict(float)
for word in allWords:
    doc = 0
    for rev in reviewsAndWords:
        if word in rev:
            doc += 1
    inverseFrequencies[word] = log(totalDocuments/doc, 10)

```

```
In [14]: for word in checkWords:
        print('IDF score for word "{}": {}'.format(word, inverseFrequencies[word]))
```

```
IDF score for word "foam": 1.1378686206869628
IDF score for word "smell": 0.5379016188648442
IDF score for word "banana": 1.6777807052660807
IDF score for word "lactic": 2.920818753952375
IDF score for word "tart": 1.8068754016455382
```

```
In [15]: for word in checkWords:
        c = reviewsAndWords[0].count(word)
        print('TF-IDF score for word "{}" in first document: {}'.format(word, c*inverseFrequencies[word]))
```

```
TF-IDF score for word "foam" in first document: 2.2757372413739256
TF-IDF score for word "smell" in first document: 0.5379016188648442
TF-IDF score for word "banana" in first document: 3.3555614105321614
TF-IDF score for word "lactic" in first document: 5.84163750790475
TF-IDF score for word "tart" in first document: 1.8068754016455382
```

```
In [16]: v1 = []
        v2 = []
        for word in allWords:
            c1 = reviewsAndWords[0].count(word)
            c2 = reviewsAndWords[1].count(word)
            v1.append(c1*inverseFrequencies[word])
            v2.append(c2*inverseFrequencies[word])
        def cos_sim(a, b):
            """Takes 2 vectors a, b and returns the cosine similarity according
            to the definition of the dot product
            """
            dot_product = np.dot(a, b)
            norm_a = np.linalg.norm(a)
            norm_b = np.linalg.norm(b)
            if norm_a == 0 or norm_b == 0:
                return 0
            return dot_product / (norm_a * norm_b)
        print('Cosine similarity of first and second review: {}'.format(cos_sim(v1, v2)))
```

```
Cosine similarity of first and second review: 0.06588193974744382
```

```
In [17]: vectors = []
        for i in range(len(reviewsAndWords)):
            v = []
            for word in allWords:
                c = reviewsAndWords[i].count(word)
                v.append(c*inverseFrequencies[word])
            vectors.append(v)
```

```

In [18]: v1 = vectors[0]

maxCos = 0
index = 1
maxIndex = -1
for v in vectors[1:]:
    cos = cos_sim(v1,v)
    if cos > maxCos:
        maxCos = cos
        maxIndex = index
    index += 1
print('Review with highest cosine similarity to first review has beer
Id "{}" and profileName "{}".format(data[maxIndex]['beer/beerId'], d
ata[maxIndex]['user/profileName']))

```

Review with highest cosine similarity to first review has beerId "72146" and profileName "spicelab"

```

In [19]: vectors = []
for i in range(len(reviewsAndWords)):
    v = []
    v.append(1)
    for word in toBeFound:
        c = reviewsAndWords[i].count(word)
        v.append(c*inverseFrequencies[word])
    vectors.append(v)
X = vectors.copy()
y = [d['review/overall'] for d in data]

clf = linear_model.Ridge(1, fit_intercept=False)
clf.fit(X, y)
theta = clf.coef_
predictions = clf.predict(X)
err = mean_squared_error(y,predictions)
print('MSE of TF-IDF model: {}'.format(err))

```

MSE of TF-IDF model: 0.27875956007772285

```

In [4]: ### Model testing
def getData(reviews, train, uni, punc, tfidf, trainWords=None):
    revWords = []
    allWords = set()
    for review in reviews:
        text = review['review/text']
        if not punc:
            words = ((text.translate(translator)).lower()).split()
        else:
            words = re.findall(r"[\w']+|[.,!?!;]", text.lower())
        revWords.append(words)
    if uni:
        revCounts = []
        wordInverseFreq = defaultdict(int)
        for rev in revWords:
            wordCounts = defaultdict(int)
            for word in rev:
                allWords.add(word)
                wordCounts[word] += 1
            revCounts.append(wordCounts)
        total = len(reviews)
        for word in allWords:
            count = 0
            for rev in revCounts:
                if word in rev:
                    count += 1
            wordInverseFreq[word] = log(total/count,10)
    else:
        revCounts = []
        wordInverseFreq = defaultdict(int)
        for rev in revWords:
            pairCounts = defaultdict(int)
            for i in range(len(rev)-1):
                pair = (rev[i], rev[i+1])
                allWords.add(pair)
                pairCounts[pair] += 1
            revCounts.append(pairCounts)
        total = len(reviews)
        for pair in allWords:
            count = 0
            for rev in revCounts:
                if pair in rev:
                    count += 1
            wordInverseFreq[pair] = log(total/count,10)
    if not tfidf:
        if train:
            X = []
            for rev in revCounts:
                feat = []
                for word in allWords:
                    if word in rev.keys():
                        feat.append(rev[word])
                    else:
                        feat.append(0)
            X.append(feat)

```

```

        else:
            X = []
            for rev in revCounts:
                feat = []
                for word in trainWords:
                    if word in rev.keys():
                        feat.append(rev[word])
                    else:
                        feat.append(0)
                X.append(feat)
    else:
        if train:
            X = []
            for rev in revCounts:
                feat = []
                for word in allWords:
                    if word in rev.keys():
                        feat.append(rev[word] * wordInverseFreq[word]
]
                                else:
                                    feat.append(0)
                                X.append(feat)
        else:
            X = []
            for rev in revCounts:
                feat = []
                for word in trainWords:
                    if word in rev.keys():
                        feat.append(rev[word] * wordInverseFreq[word]
]
                                else:
                                    feat.append(0)
                                X.append(feat)
    return X, list(allWords)

```

```

In [7]: def train_and_test(unigram,punctuation,tfidf):
        X_train, words = getData(trainData, True, unigram,punctuation,tfi
df)
        y_train = [r['review/overall'] for r in trainData]
        X_val, _ = getData(valData,False,unigram,punctuation,tfidf,words)
        y_val = [r['review/overall'] for r in valData]
        regularizers = [0.01,0.1,1,10,100]
        for reg in regularizers:
            clf = linear_model.Ridge(reg, fit_intercept=False)
            clf.fit(X_train, y_train)
            theta = clf.coef_
            predictions = clf.predict(X_val)
            err = mean_squared_error(y_val,predictions)
            print('Words: {}, Punctuation: {}, Model: {}, Regularizer: {}
, MSE: {}'.format('Unigram' if unigram else 'Bigram','Kept' if punctu
ation else 'Removed','TF-IDF' if tfidf else 'Counts',reg,err))

```



```
In [8]: unigram = [False, True]
punctuation = [True, False]
tfidf = [True, False]

from multiprocessing import Process

processes = []

for j in punctuation:
    for k in tfidf:
        train_and_test(False, j, k)

for j in punctuation:
    for k in tfidf:
        p = Process(target=train_and_test, args=((True, j, k)))
        p.start()
        p.join()
        processes.append(p)
for p in processes:
    p.join()
```

Words: Bigram, Punctuation: Kept, Model: TF-IDF, Regularizer: 0.01, MSE: 1.8862181310759654  
Words: Bigram, Punctuation: Kept, Model: TF-IDF, Regularizer: 0.1, MSE: 1.8861658649250537  
Words: Bigram, Punctuation: Kept, Model: TF-IDF, Regularizer: 1, MSE: 1.885647097566243  
Words: Bigram, Punctuation: Kept, Model: TF-IDF, Regularizer: 10, MSE: 1.8808223634720744  
Words: Bigram, Punctuation: Kept, Model: TF-IDF, Regularizer: 100, MSE: 1.8548149170135138  
Words: Bigram, Punctuation: Kept, Model: Counts, Regularizer: 0.01, MSE: 2.0245113627375444  
Words: Bigram, Punctuation: Kept, Model: Counts, Regularizer: 0.1, MSE: 2.023304775836114  
Words: Bigram, Punctuation: Kept, Model: Counts, Regularizer: 1, MSE: 2.0118209329540218  
Words: Bigram, Punctuation: Kept, Model: Counts, Regularizer: 10, MSE: 1.933035297458138  
Words: Bigram, Punctuation: Kept, Model: Counts, Regularizer: 100, MSE: 1.794753905512156  
Words: Bigram, Punctuation: Removed, Model: TF-IDF, Regularizer: 0.01, MSE: 2.167417092541125  
Words: Bigram, Punctuation: Removed, Model: TF-IDF, Regularizer: 0.1, MSE: 2.167395081397154  
Words: Bigram, Punctuation: Removed, Model: TF-IDF, Regularizer: 1, MSE: 2.167177458626135  
Words: Bigram, Punctuation: Removed, Model: TF-IDF, Regularizer: 10, MSE: 2.165233997268316  
Words: Bigram, Punctuation: Removed, Model: TF-IDF, Regularizer: 100, MSE: 2.1604211159781204  
Words: Bigram, Punctuation: Removed, Model: Counts, Regularizer: 0.01, MSE: 2.3031320540739304  
Words: Bigram, Punctuation: Removed, Model: Counts, Regularizer: 0.1, MSE: 2.302177299615284  
Words: Bigram, Punctuation: Removed, Model: Counts, Regularizer: 1, MSE: 2.2931072880386636  
Words: Bigram, Punctuation: Removed, Model: Counts, Regularizer: 10, MSE: 2.2325865651085564  
Words: Bigram, Punctuation: Removed, Model: Counts, Regularizer: 100, MSE: 2.1809703271148106  
Words: Unigram, Punctuation: Kept, Model: TF-IDF, Regularizer: 0.01, MSE: 4.01055826797125  
Words: Unigram, Punctuation: Kept, Model: TF-IDF, Regularizer: 0.1, MSE: 3.948424080348474  
Words: Unigram, Punctuation: Kept, Model: TF-IDF, Regularizer: 1, MSE: 3.5364562233438903  
Words: Unigram, Punctuation: Kept, Model: TF-IDF, Regularizer: 10, MSE: 2.568415049160528  
Words: Unigram, Punctuation: Kept, Model: TF-IDF, Regularizer: 100, MSE: 1.8861795490128672  
Words: Unigram, Punctuation: Kept, Model: Counts, Regularizer: 0.01, MSE: 4.322541025784355  
Words: Unigram, Punctuation: Kept, Model: Counts, Regularizer: 0.1, MSE: 3.891389329026427  
Words: Unigram, Punctuation: Kept, Model: Counts, Regularizer: 1, MSE: 2.7900857443768685

Words: Unigram, Punctuation: Kept, Model: Counts, Regularizer: 10, MSE: 1.9151691017847603  
 Words: Unigram, Punctuation: Kept, Model: Counts, Regularizer: 100, MSE: 1.607167223096318  
 Words: Unigram, Punctuation: Removed, Model: TF-IDF, Regularizer: 0.01, MSE: 3.4711037409717  
 Words: Unigram, Punctuation: Removed, Model: TF-IDF, Regularizer: 0.1, MSE: 3.4379850742336084  
 Words: Unigram, Punctuation: Removed, Model: TF-IDF, Regularizer: 1, MSE: 3.196840379015003  
 Words: Unigram, Punctuation: Removed, Model: TF-IDF, Regularizer: 10, MSE: 2.495611933195827  
 Words: Unigram, Punctuation: Removed, Model: TF-IDF, Regularizer: 100, MSE: 1.9102693980827168  
 Words: Unigram, Punctuation: Removed, Model: Counts, Regularizer: 0.01, MSE: 3.786931886877045  
 Words: Unigram, Punctuation: Removed, Model: Counts, Regularizer: 0.1, MSE: 3.534393088459552  
 Words: Unigram, Punctuation: Removed, Model: Counts, Regularizer: 1, MSE: 2.7259494347253814  
 Words: Unigram, Punctuation: Removed, Model: Counts, Regularizer: 10, MSE: 1.94142584543845  
 Words: Unigram, Punctuation: Removed, Model: Counts, Regularizer: 100, MSE: 1.6493934955017735

```
In [9]: ### Best model's performance on the test set
X_train, words = getData(trainData, True, True, True, False)
y_train = [r['review/overall'] for r in trainData]
X_val, _ = getData(testData, False, True, True, False, words)
y_val = [r['review/overall'] for r in testData]
clf = linear_model.Ridge(100, fit_intercept=False)
clf.fit(X_train, y_train)
theta = clf.coef_
predictions = clf.predict(X_val)
err = mean_squared_error(y_val, predictions)
print('Words: {}, Punctuation: {}, Model: {}, Regularizer: {}, MSE: {}'
      .format('Unigram', 'Kept', 'Counts', 100, err))
```

Words: Unigram, Punctuation: Kept, Model: Counts, Regularizer: 100, MSE: 1.6688250630614938